

Playing Atari with Deep Reinforcement Learning

Ha-larm

Abstract

- First deep learning model using reinforcement learning
 - Successfully learn control policies directly
 - From high-dimensional sensory input(pixels)
- CNN model trained on a variant of Q-learning
 - input:raw pixel,output:a value function estimation future reward
- Applied seven Atari 2600 games (no adjustment)
 - Outperforms ALL previous approaches on six games
 - Surpasses a human expert on three games


Problem to Solve

- Learning directly from high dimensional sensory input
 - Long-standing challenges of RL
- Most successful RL relies on hand-crafted features

Solution

- Most DL requires hand labeled training data
 - RL must learn from a scalar reward signal
 - Reward signal is often sparse, noisy, and delayed
 - delay between actions and resulting rewards can be thousand time steps
 - CNN with a variant Q-learning
- Most DL assumes data samples are independent
 - RL encounters sequences of highly correlated states
 - Experience replay

Agent and Environment



- ▶ At each step t the agent:
 - ▶ Executes action a_t
 - ▶ Receives observation o_t
 - ▶ Receives scalar reward r_t
- ▶ The environment:
 - ▶ Receives action a_t
 - ▶ Emits observation o_{t+1}
 - ▶ Emits scalar reward r_{t+1}

State



- Experience is a sequence of observations, actions, rewards

$$o_1, r_1, a_1, \dots, a_{t-1}, o_t, r_t$$

- The **state** is a summary of experience

$$s_t = f(o_1, r_1, a_1, \dots, a_{t-1}, o_t, r_t)$$

- In a fully observed environment

$$s_t = f(o_t)$$

Policy

- Policy is the agent's behaviour
- Deterministic policy: $a = \pi(s)$
- Stochastic policy: $\pi(a|s) = P[a|s]$ (확률)

Value Function

- ▶ A **value function** is a prediction of future reward
 - ▶ “How much reward will I get from action a in state s ?”
- ▶ **Q -value function** gives expected total reward
 - ▶ from state s and action a
 - ▶ under policy π
 - ▶ with discount factor γ

$$Q^{\pi}(s, a) = \mathbb{E} [r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots \mid s, a]$$

- ▶ Value functions decompose into a Bellman equation

$$Q^{\pi}(s, a) = \mathbb{E}_{s', a'} [r + \gamma Q^{\pi}(s', a') \mid s, a]$$

Optimal Value Functions

- ▶ An optimal value function is the maximum achievable value

$$Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a) = Q^{\pi^*}(s, a)$$

- ▶ Once we have Q^* we can act optimally,

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$

- ▶ Optimal value maximises over all decisions. Informally:

$$\begin{aligned} Q^*(s, a) &= r_{t+1} + \gamma \max_{a_{t+1}} r_{t+2} + \gamma^2 \max_{a_{t+2}} r_{t+3} + \dots \\ &= r_{t+1} + \gamma \max_{a_{t+1}} Q^*(s_{t+1}, a_{t+1}) \end{aligned}$$


- ▶ Formally, optimal values decompose into a Bellman equation

$$Q^*(s, a) = \mathbb{E}_{s'} \left[r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]$$

Q-Networks

Represent value function by **Q-network** with weights **w**

$$Q(s, a, \mathbf{w}) \approx Q^*(s, a)$$



Q-Learning

- ▶ Optimal Q-values should obey Bellman equation

$$Q^*(s, a) = \mathbb{E}_{s'} \left[r + \gamma \max_{a'} Q(s', a')^* \mid s, a \right]$$

- ▶ Treat right-hand side $r + \gamma \max_{a'} Q(s', a', \mathbf{w})$ as a target
- ▶ Minimise MSE loss by stochastic gradient descent


$$l = \left(\underbrace{r + \gamma \max_{a'} Q(s', a', \mathbf{w})}_{\text{target}} - \underbrace{Q(s, a, \mathbf{w})}_{\text{current}} \right)^2$$

Converge

- Converges to max Q using table lookup representation
- But **diverges** using NN due to:
 - Correlations between samples
 - Non-stationary targets

Deep Q-Networks (DQN): Experience Replay

To remove correlations, build data-set from agent's own experience



Sample experiences from data-set and apply update




Figure 2: The two plots on the left show average reward per episode on Breakout and Seaquest respectively during training. The statistics were computed by running an ϵ -greedy policy with $\epsilon = 0.05$ for 10000 steps. The two plots on the right show the average maximum predicted action-value of a held out set of states on Breakout and Seaquest respectively. One epoch corresponds to 50000 minibatch weight updates or roughly 30 minutes of training time.

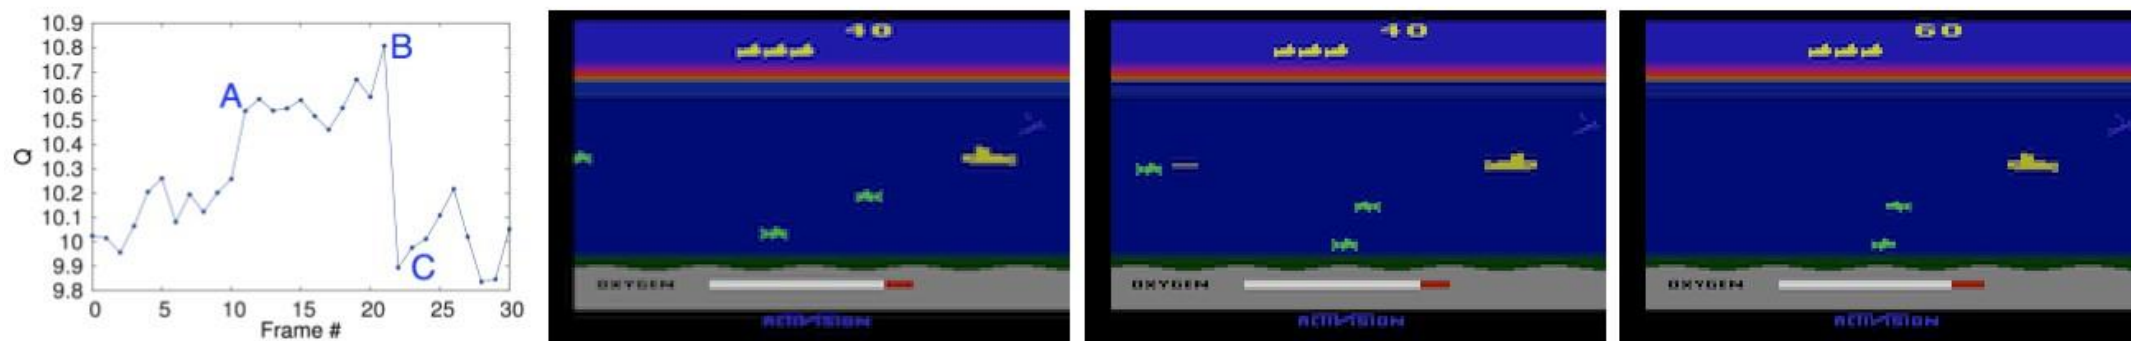


Figure 3: The leftmost plot shows the predicted value function for a 30 frame segment of the game Seaquest. The three screenshots correspond to the frames labeled by A, B, and C respectively.

Q & A

Thank you!

Paper:PlayingAtariwithDeepReinforcementLearning